

OFF-ROAD DESERT SEMANTIC SEGMENTATION

Pixel-wise terrain labelling for synthetic off-road desert scenes

using ResNet50 + Feature Pyramid Network Decoder

Team AIDHUNIK · Hack For Green Bharat · February 26, 2026

The Problem

Why terrain perception matters for autonomous off-road systems



Harsh Terrain Diversity

10 semantically distinct terrain classes
that look visually similar in desert
environments make classification hard.



Sim-to-Real Gap

Synthetic data has distribution shifts vs.
real imagery. Models must generalise
across lighting, texture, and geometry.



Class Imbalance

Sky and ground dominate pixel counts;
rare classes like bushes/rocks are under-
represented, biasing naïve models.

Dataset

Duality AI — Synthetic Desert Off-Road Scenes

2 857

Training
images

317

Validation
images

10

Semantic
classes

476×266

Image
resolution

Mask pixel label mapping:

Background (raw=0)

Trees (raw=100)

Dry Grass (raw=200)

Rocks (raw=300)

Bushes (raw=500)

Sky (raw=550)

Ground Clutter (raw=700)

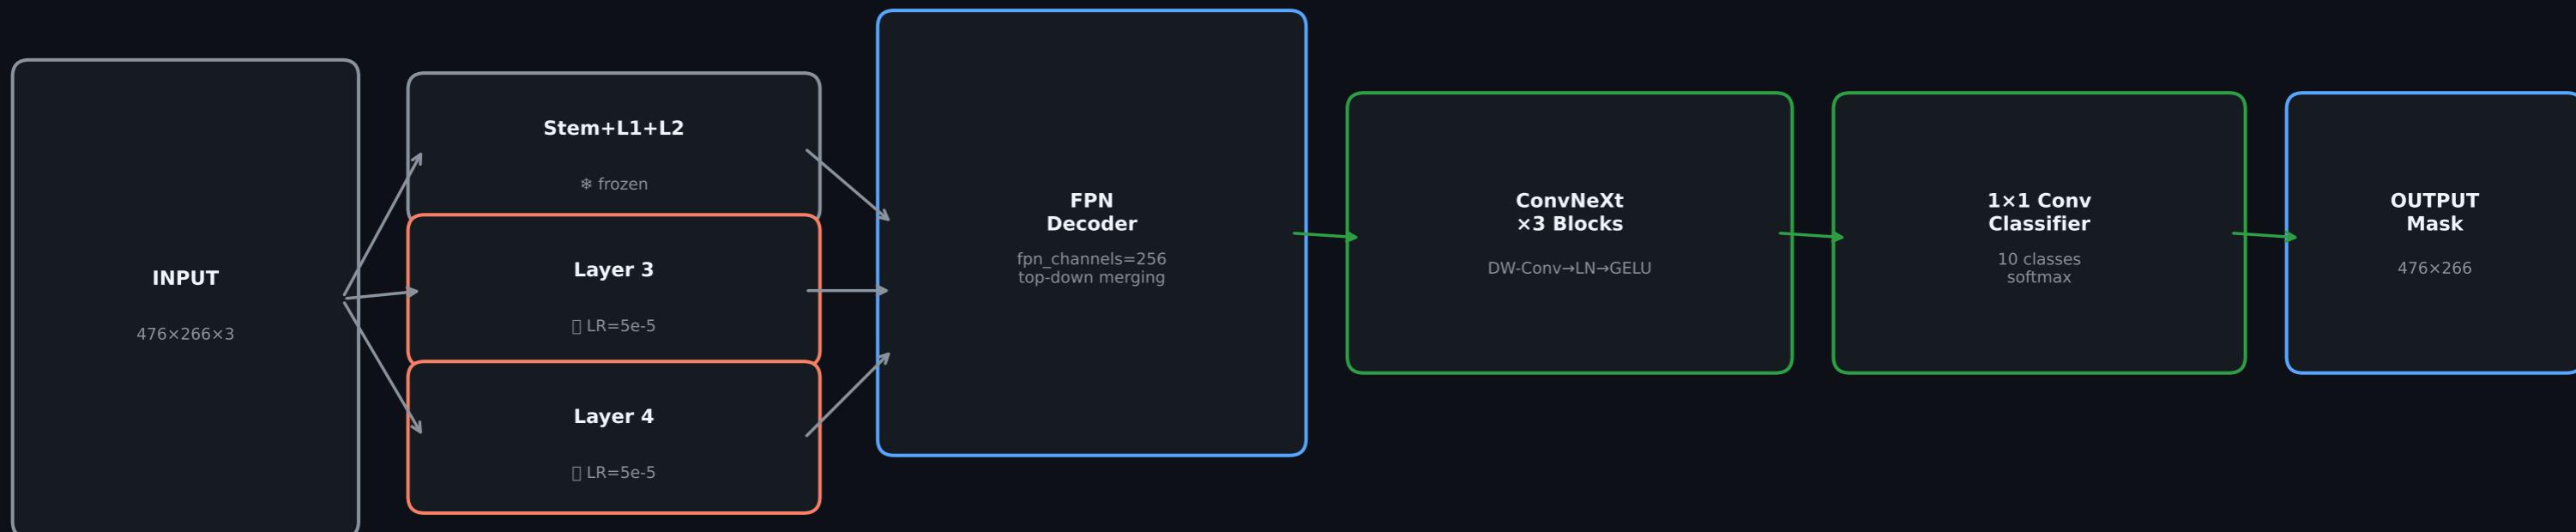
Terrain (raw=800)

Shrubs (raw=7100)

Trail (raw=10000)

Architecture

ResNet50 Multi-Scale Backbone + FPN Decoder + ConvNeXt Head

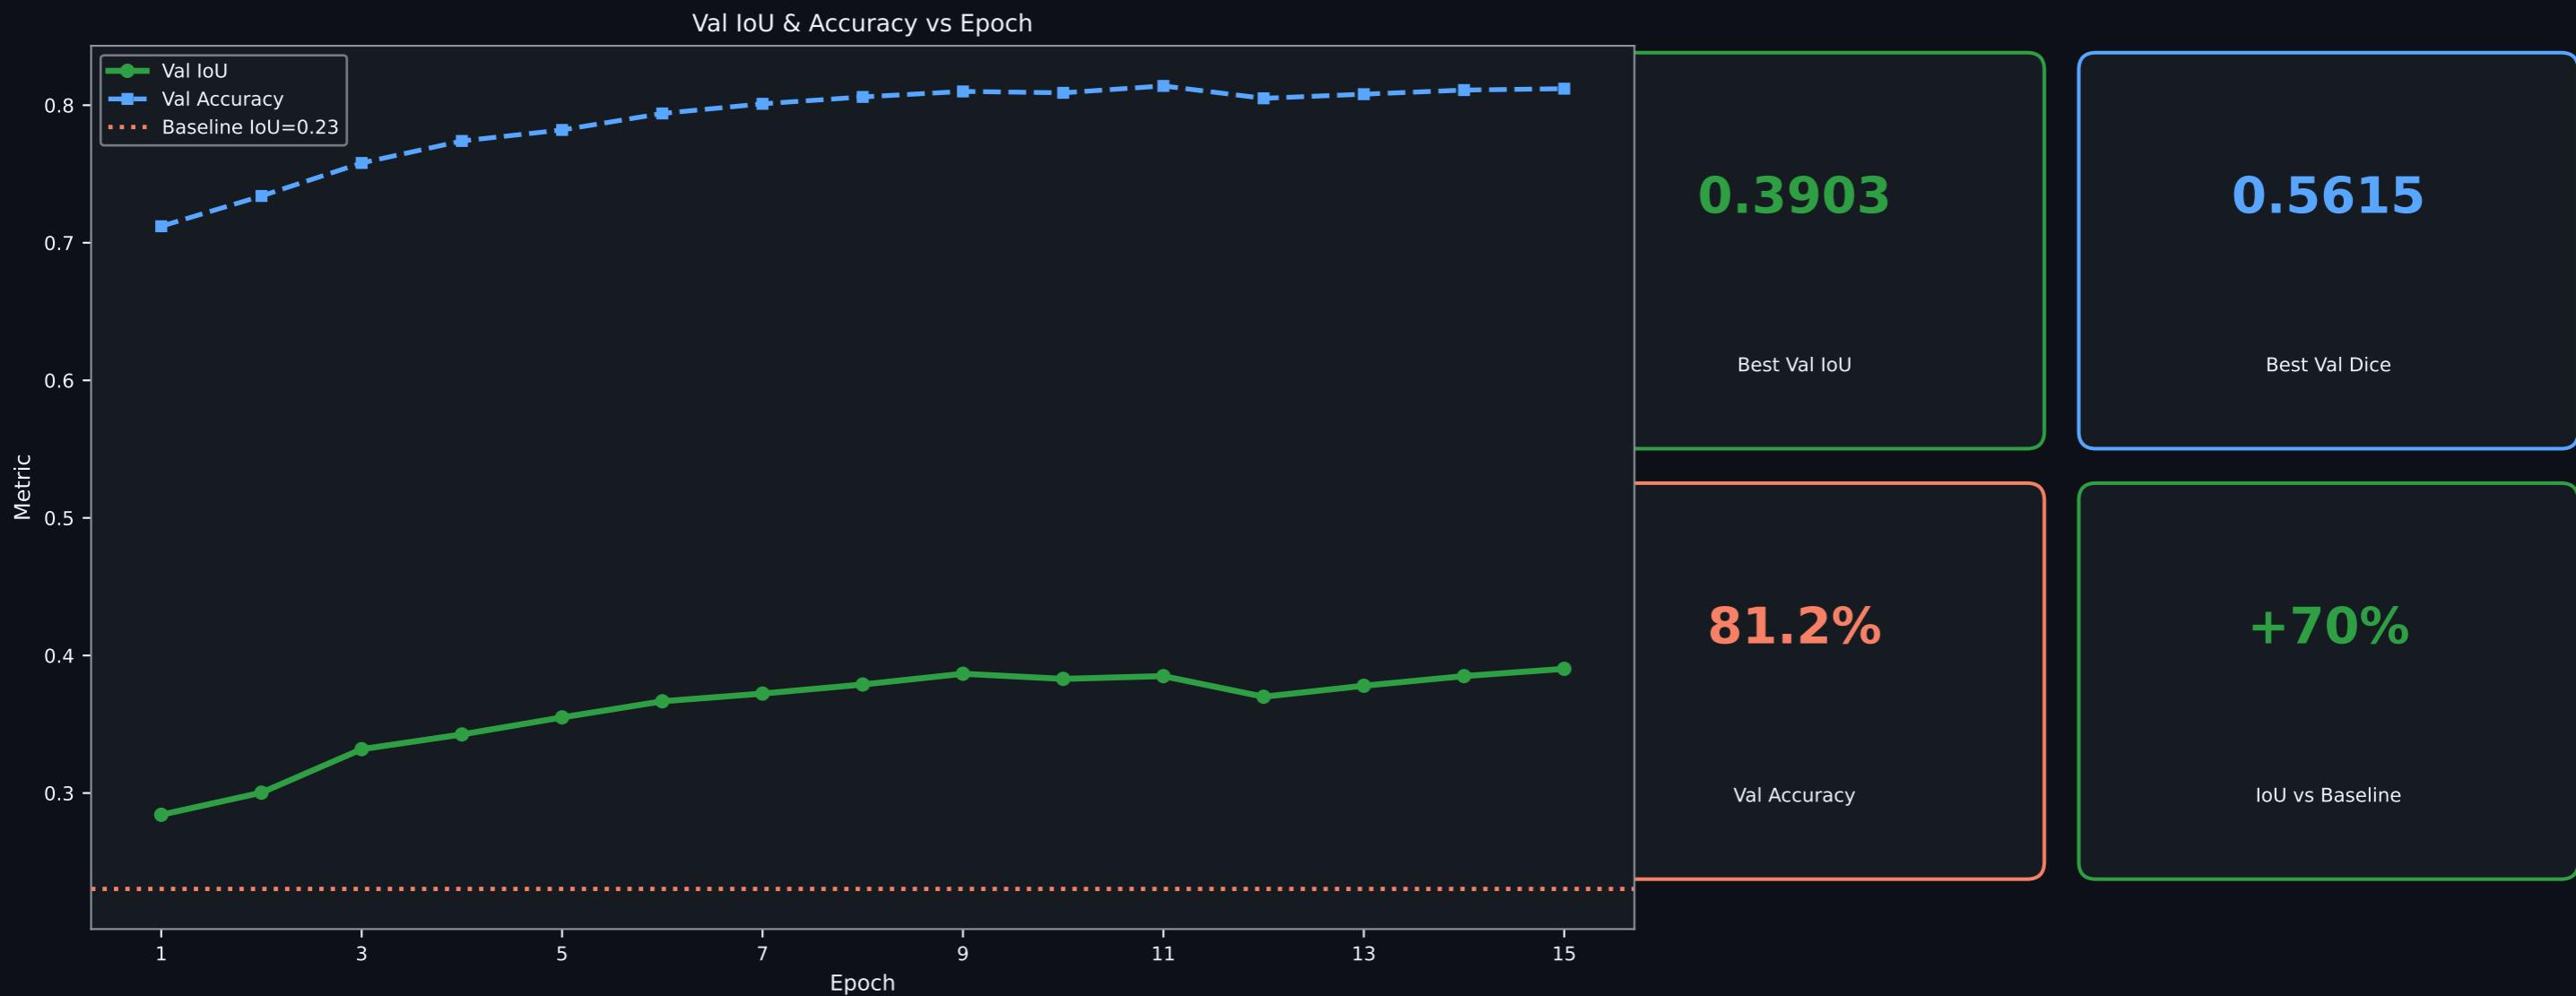


$$\text{Loss} = \text{CrossEntropy(class-weighted, label-smooth=0.05)} + 0.5 \times \text{SoftDice}$$

Optimiser: AdamW · Scheduler: CosineAnnealingLR $T_0=10$ · Head LR=5e-4 · Backbone LR=5e-5

Training Results

V3 FPN — 15 epochs on single T4 GPU



Model Comparison

Baseline v1 → V3 FPN

Model	Epochs	Val IoU	Val Dice	Val Accuracy	Val Loss
Baseline v1 (frozen backbone)	10	0.2305	0.3633	66.08%	0.96
V3 FPN ★ (partial fine-tune)	15	0.3903	0.5615	81.20%	1.79
Δ improvement	—	+69.6%	+54.5%	+22.9 pp	—

Key Innovations

What made V3 outperform the baseline by 70%

1

Multi-Scale FPN

Fuse features from layer2, layer3, layer4
Captures both fine detail and coarse semantics
Top-down pathway with lateral skip connections

2

Partial Backbone Fine-Tuning

Layer3 + Layer4 unfrozen at LR=5e-5
Adapts ImageNet weights to desert textures
10x lower LR than head to avoid forgetting

3

Combined Loss + Class Weights

CE (label-smooth=0.05) + 0.5 × SoftDice
Inverse-frequency class weights tackle imbalance
Dice loss directly optimises the IoU metric

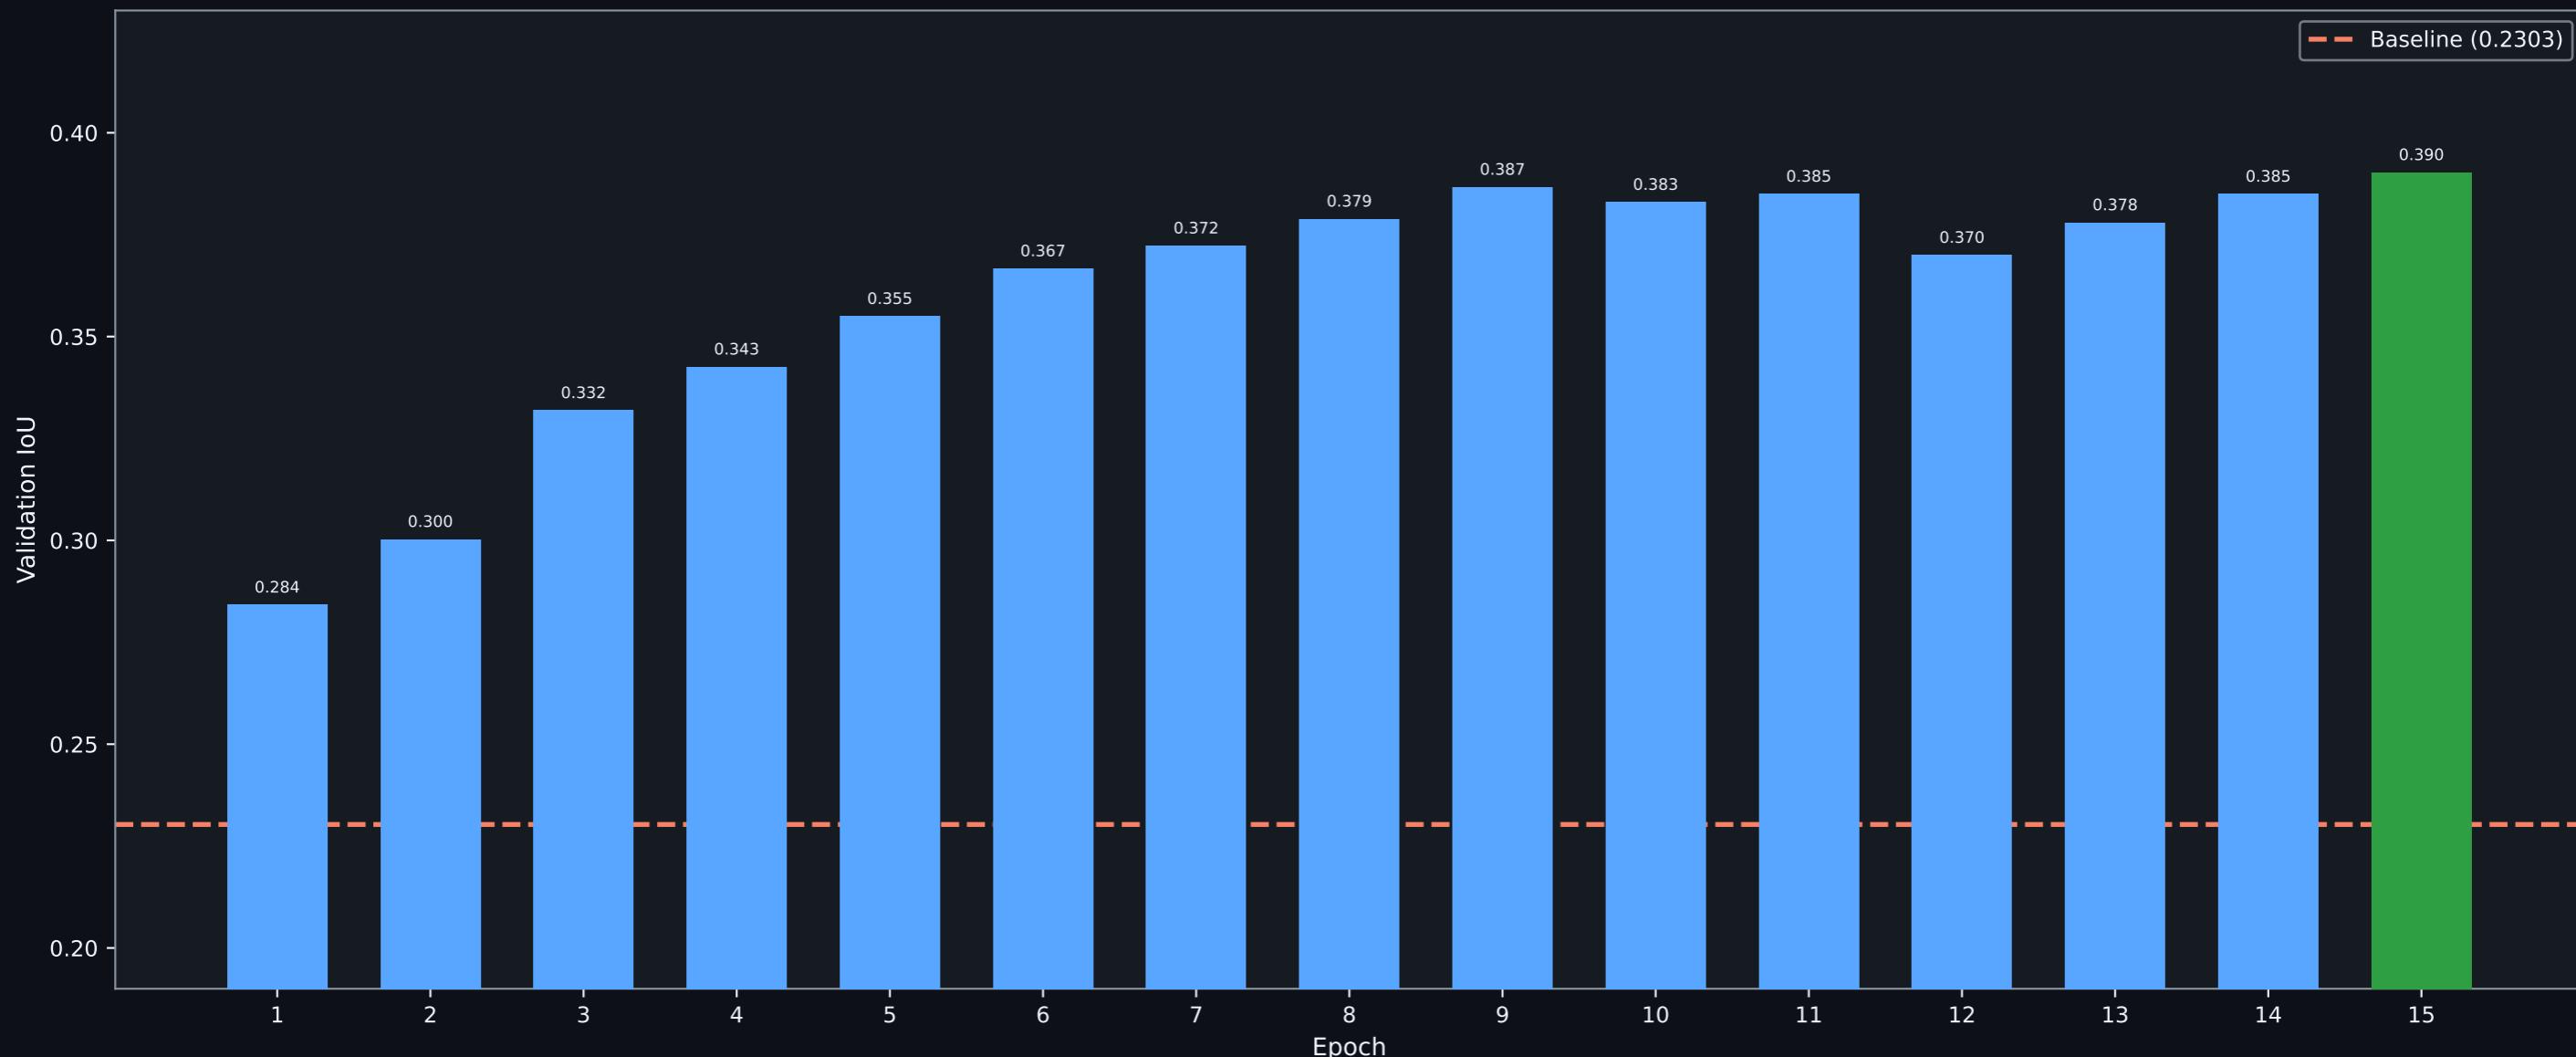
4

ConvNeXt-Style Head

3 × ConvNeXt blocks: DW-Conv → LN → GELU
More expressive than plain conv stacks
CosineAnnealing LR restarts every $T_0=10$ epochs

Per-Epoch Val IoU

Training progression over 15 epochs



Best: Epoch 15 → Val IoU = 0.3903 (+69.6% over baseline)

Submission Package

What we're delivering

□ **SUBMISSION/src/**

train_segmentation.py · test_segmentation.py · visualize.py

□ **models/**

segmentation_head_best.pth (120 MB — best checkpoint, epoch 15)

□ **train_stats/v3/**

evaluation_metrics_v3.txt · 4 training curve plots (PNG)

□ **README.md & REPORT.md**

Architecture docs, results table, setup & usage instructions

□ **requirements.txt**

torch · torchvision · numpy · opencv-python · albumentations · tqdm

□ **GitHub Repo**

github.com/Akum030/duality-desert-segmentation

Thank You!

Team AIDHUNIK

0.3903

Val IoU

0.5615

Val Dice

81.2%

Val Acc

+70%

Δ IoU

github.com/Akum030/duality-desert-segmentation

Hack For Green Bharat · February 26, 2026